

A Modular Multi-Omics Data Integration Pipeline for Precision Oncology: Design, Implementation, and Validation on TCGA-BRCA

Nidhal Zitouni
Faculty of Sciences of Monastir
Tunisia
nidhal.zitouni@fsm.tn

Mohsen Maraoui
Faculty of Sciences of Monastir
Tunisia
mohsen.maraoui@fsm.tn

Abstract

Integrating heterogeneous molecular datasets remains a formidable challenge in translational oncology, where high-dimensional noise and cross-modal inconsistencies often hinder biomarker discovery. We present MultiOmicsPipeline, a modular Python framework designed to orchestrate the complete lifecycle of multi-omics integration. Our architecture employs a five-phase directed acyclic graph (DAG) encompassing automated data ingestion, robust imputation via k-Nearest Neighbors (k-NN), multi-strategy normalization, PCA-based latent feature extraction, and FHIR R4-compliant clinical reporting. Validated on the TCGA-BRCA cohort ($n = 1,098$), the pipeline successfully harmonized transcriptomic, epigenomic, and clinical layers, achieving a composite data quality score of 0.989 ± 0.004 across six analytical dimensions. By reducing end-to-end processing time to 68.4 minutes on standard hardware while ensuring FAIR-compliant interoperability, this framework provides a scalable and reproducible foundation for large-scale precision medicine initiatives.

CCS Concepts

• Applied computing → Bioinformatics.

Keywords

Multi-omics Integration, PCA, Precision Oncology, FHIR R4, TCGA-BRCA, Pipeline Orchestration, Bioinformatics Interoperability

1 Introduction

Precision oncology is fundamentally predicated on the synergistic analysis of genomic, transcriptomic, and epigenomic landscapes to uncover clinically actionable biomarkers [1]. However, the integration of these disparate molecular layers is complicated by significant technical barriers. Researchers must reconcile massive dimensionality differentials (ranging from 10^3 to 10^6 features), mitigate batch effects inherent in multi-center sequencing [20], and manage incomplete sample overlap, all while ensuring the final output is structured for clinical downstream utility.

Existing tools, such as TCGAbiolinks [7], cBioPortal, and Galaxy-P, offer robust capabilities for specific analytical tasks but often lack the end-to-end cohesion required for rapid, reproducible pipeline execution. In practice, bioinformaticians are frequently forced to develop ad-hoc scripts to bridge the gaps between ingestion, imputation, and latent-factor integration, leading to technical debt and reproducibility crises.

To address this fragmentation, we introduce MultiOmicsPipeline, a unified Python-based orchestration framework¹. Our primary contributions include: (1) a modular preprocessing engine featuring scalable, density-aware imputation; (2) a robust PCA-based intermediate integration strategy that minimizes high-dimensional noise prior to fusion; (3) a rigorous six-dimension quality assessment framework with bootstrap validation; and (4) an automated export module compliant with FHIR R4 Genomics Reporting standards [21], bridging the gap between bioinformatics research and clinical informatics systems.

2 Background and Related Work

2.1 Multi-Omics Data Taxonomy

Table 1 summarizes the principal molecular layers typically encountered in cohorts like TCGA-BRCA.

Table 1: Taxonomy of molecular layers in multi-omics cohorts.

Layer	Technology	Metric	Dimensions
Genomics	WGS / WES	SNVs, CNVs	3×10^6
Transcriptomics	RNA-Seq	Gene expression	6×10^4
Epigenomics	450K / 850K	DNA Methylation	4.8×10^5
Proteomics	RPPA / LC-MS	Protein abundance	2×10^4
Metabolomics	LC-MS / NMR	Small molecules	5×10^3

2.2 Integration Paradigms

Multi-omics integration generally falls into three categories. **Early fusion** involves the direct concatenation of high-dimensional feature matrices:

$$\mathbf{X}_{\text{int}} = [\mathbf{X}^{(1)} \parallel \mathbf{X}^{(2)} \parallel \dots \parallel \mathbf{X}^{(M)}] \quad (1)$$

Intermediate fusion employs dimensionality reduction to project each modality into a unified latent space prior to integration:

$$\mathbf{Z}^{(m)} = \phi_m(\mathbf{x}^{(m)}), \quad \mathbf{X}_{\text{integrated}} = \mathcal{G}(\mathbf{Z}^{(1)}, \dots, \mathbf{Z}^{(M)}) \quad (2)$$

Late fusion aggregates the decision-level outputs of modality-specific predictive models:

$$\hat{y} = \mathcal{F}(f_1(\mathbf{X}^{(1)}), f_2(\mathbf{X}^{(2)}), \dots, f_M(\mathbf{X}^{(M)})) \quad (3)$$

¹<https://github.com/ZitouniNidhal/Multi-Omics-Clinical-Integration-Pipeline-MOCIP>

2.3 Identified Gaps

Despite the availability of excellent standalone tools, configuring a reproducible pipeline from raw data to clinical format requires significant manual effort. Table 2 outlines how our pipeline addresses these specific workflow gaps.

Table 2: Workflow bottlenecks and MultiOmicsPipeline solutions.

Challenge	Impact	Our Solution
Data Silos	Lack of clinical context	Built-in FHIR R4 mapping
Validation Gap	Unreliable comparisons	6D quality scoring engine
Code Rigidity	High adaptation cost	Hexagonal + DAG design
Fragmentation	Extended prep time	5-stage orchestration

3 Pipeline Architecture

3.1 Design Principles

The development of MultiOmicsPipeline was guided by four fundamental architectural paradigms: *Modular Extensibility*, allowing for the seamless hot-swapping of analytical components; *Deterministic Reproducibility*, achieved through version-controlled YAML configuration schemas; *Comprehensive Observability*, ensuring telemetry and logging at every pipeline junction; and *Semantic Interoperability*, prioritizing FAIR data principles and standardized clinical ontologies.

3.2 Hexagonal Orchestration

We implemented a five-tier hexagonal architecture to decouple core domain logic from external data sources and interfaces (Figure 1). This structural separation ensures that the integration logic remains agnostic to specific sequencing platforms or database technologies.

- **L0 (Data Ingress):** Manages asynchronous API interactions with the Genomic Data Commons (GDC) and local repository mounting.
- **L1 (Transformation):** Executes vectorized imputation and normalization workflows.
- **L2 (Synergy):** Performs cross-modal sample alignment and latent feature extraction via PCA.
- **L3 (Analytics):** Orchestrates predictive modeling and iterative validation logic.
- **L4 (Presentation):** Exposes functionality through a RESTful API and a Command Line Interface (CLI).

3.3 Implementation: Core Classes

The execution flow is coordinated by the MultiOmicsPipeline class, which abstracts the complexity of the underlying steps:

Listing 1: Simplified pipeline orchestration snippet.

```
1 class MultiOmicsPipeline:
2     def run(self, omic_path, clinical_path, output_dir):
3         omic = self.load_data(omic_path)
4         clinical = self.load_data(clinical_path)
5         processed = self.preprocess_data(omic, clinical)
```

```
6         integrated = self.integrate_data(processed)
7         ml_data, results = self.run_model(integrated)
8         self.export_data(integrated, output_dir)
```

Specific functionalities are delegated to dedicated modules, such as MissingValueHandler for KNN imputation, OmicsNormalizer for method-specific scaling (e.g., TPM, DESeq2), SampleAlignment for matching patient IDs, and MultiOmicsFusion for managing the different integration strategies.

3.4 Technology Stack

Table 3: The MultiOmicsPipeline technology stack.

Layer	Technologies	Purpose
Business	Python 3.10, Pydantic	Core logic and schema validation
Processing	Pandas, NumPy, Scikit-learn	Vectorized manipulation and ML
Integration	Dask, Scikit-learn	Parallelized latent-factor analysis
Storage	Parquet, PostgreSQL	High-performance columnar I/O
Infrastructure	Docker, Kubernetes	Scalable containerized deployment

4 Dataset: TCGA-BRCA Validation

To evaluate the pipeline’s robustness, we utilized the Breast Invasive Carcinoma (TCGA-BRCA) cohort [19], which provides a complex landscape of multi-modal data. Our validation focused on a subset of 1,098 primary tumor samples with overlapping omics signatures, as detailed in Table 4. This cohort presents significant challenges, including an 8.7% mean missingness rate across RNA-Seq and Methylation layers, making it an ideal benchmark for our imputation and alignment modules.

Table 4: Characteristics of the TCGA-BRCA validation cohort.

Attribute	Value
Samples (Primary Tumors)	1,098 Patients
Omics Layers	RNA-Seq, 450K, WES
Clinical Variables	45 per Patient
Raw Data Volume	1.27 GB
Baseline Quality Score	0.912 / 1.000
Initial Missingness	8.7% (Mean)

5 Data Preprocessing Strategies

5.1 Adaptive Missing Value Imputation

High-throughput omics experiments are frequently plagued by missing data due to technical limitations or biological dropout. Our MissingValueHandler implements a density-aware k-Nearest




Figure 1: Architectural Workflow: From raw multi-modal sequencing data to interoperable clinical reporting. The pipeline is structured into four distinct phases ensuring reproducibility and rigorous standardization.

Neighbors (k-NN) strategy [11], which leverages local feature similarities to estimate missing values without distorting the underlying data distribution. For computational efficiency on large cohorts, the handler falls back to median-based imputation for low-variance features while proactively pruning features with missingness exceeding a configurable threshold (default: 20%).

5.2 Multi-Modal Normalization

Given the disparate dynamic ranges across transcriptomic and epigenomic layers, normalization is critical for preventing feature dominance. We provide a tiered normalization engine:

- **RNA-Seq:** Support for TPM (Transcripts Per Million) and DESeq2-based size factor scaling to account for library depth.
- **Proteomics/Clinical:** Robust Z-score standardization and Pareto scaling to mitigate the influence of outliers.
- **Epigenomics:** Quantile normalization to ensure consistency across methylation arrays.

6 Multi-Omics Integration Strategies

6.1 Cross-Modal Sample Alignment

A primary bottleneck in multi-omics workflows is the reconciliation of sample identifiers across disparate platforms. Our `SampleAlignment` module implements a robust matching algorithm that extracts core TCGA barcodes while supporting fuzzy matching for mislabeled clinical entries. As demonstrated in Table 5, we achieved high retention rates (> 99%) across all modalities, ensuring that the biological signal from the 1,098-patient cohort was preserved during the intersection process.

6.2 Fusion Paradigms

The framework supports three distinct fusion strategies to accommodate varying analytical objectives:

- (1) **Early Fusion (Feature-Level):** Horizontal concatenation of raw matrices, optimized with prefix-aware indexing to prevent feature collision.

Table 5: Cross-modal sample alignment and retention rates.

Modality	Initial	Aligned	Retention
RNA-Seq	1,098	1,095	99.7%
Clinical Data	1,102	1,095	99.4%
WES (Mutations)	978	975	99.7%
Methylation	894	892	99.8%
Integrated	—	1,095	99.7%

- (2) **Intermediate Fusion (Latent-Level):** Applying Principal Component Analysis (PCA) [22] to each modality independently to project high-dimensional data into a lower-dimensional latent space. This approach effectively filters technical noise while capturing the most significant biological variance.
- (3) **Late Fusion (Decision-Level):** Aggregating predictions from modality-specific models, suitable for cases with extreme modality heterogeneity.

6.3 PCA-based Dimensionality Reduction

To mitigate the "curse of dimensionality," we utilized PCA for intermediate integration. By selecting the top components that explain 95% of the variance in each layer, we reduced the feature space of the transcriptomic data from 60,000+ dimensions to approximately 350 latent factors, significantly improving the computational efficiency of downstream machine learning models without sacrificing predictive accuracy.

7 Export and Interoperability

To ensure the processed data is usable by other tools and clinical systems, we built four separate exporter classes. The pipeline can output tab-separated CSVs, versioned JSON files, Parquet files for fast I/O, and FHIR R4 Genomics Reporting formats.

Table 6: Interoperability metrics across supported export formats.

Metric	Parquet	cBioPortal	Xena	FHIR R4
Footprint (GB)	2.8	3.1	3.0	4.2
Latence (s)	2.3	4.5	3.8	5.2
Compatibility	95%	100%	98%	85%
Reproducibility	100%	95%	92%	98%

The FHIR export specifically maps the data to Patient, Observation, and DiagnosticReport resources, passing R4 conformance checks without errors.

8 Experimental Results

8.1 Quantifying Pipeline Performance

The end-to-end evaluation on the TCGA-BRCA cohort revealed substantial gains in data integrity and analytical readiness. Our results (Table 7) indicate that the preprocessing workflow effectively eliminated missingness while enhancing the Signal-to-Noise Ratio (SNR) by over 280%. Crucially, the pipeline maintained 100% reproducibility through its configuration-driven execution model.

Table 7: End-to-end pipeline performance on TCGA-BRCA.

Metric	Baseline	Processed	Integrated	Δ
Quality Score	0.912	0.987	0.991	+8.7%
Missingness (%)	8.7	0.0	0.0	−100%
SNR (dB)	3.2	8.7	12.3	+284%
Alignment (%)	97.1	99.5	99.7	+2.7 pp
Wall Time (min)	—	45.2	68.4	—

8.2 Multidimensional Quality Benchmarking

To ensure the integrated data meets the rigorous standards of clinical research, we applied a custom six-dimension quality framework (Table 8). The composite quality score of 0.989 (95% CI: [0.985, 0.993]), derived via 1,000 bootstrap repetitions, confirms that the MultiOmicsPipeline produces highly consistent and valid datasets suitable for biomarker discovery and clinical modeling.

8.3 Computational Scalability and Efficiency

Computational efficiency is a critical requirement for pan-cancer analyses. Our pipeline demonstrated near-linear $O(n)$ scaling, processing the entire 1,098-patient dataset (RNA-Seq, Methylation, and Clinical layers) in 68.4 minutes. This performance profile is particularly noteworthy given the complexity of the k-NN imputation and PCA steps, suggesting that the framework is well-suited for larger prospective cohorts.

Table 8: Multidimensional quality assessment of the integrated output.

Dimension	Score	95% CI	Methodology
Completeness	1.000	[1.000, 1.000]	Exhaustive null check
Accuracy	0.982	[0.975, 0.989]	Ground-truth overlap
Consistency	0.965	[0.955, 0.975]	Internal variance check
Validity	0.991	[0.985, 0.997]	Clinical domain rules
Uniqueness	1.000	[1.000, 1.000]	Primary key collision
Timeliness	0.998	[0.995, 1.000]	Metadata freshness
Composite	0.989	[0.985, 0.993]	Bootstrap (1k rep.)

9 Discussion and Future Horizons

9.1 Methodological Implications

Our findings underscore the importance of automated orchestration in multi-omics research. The flexibility afforded by the YAML-driven configuration allows researchers to pivot between imputation strategies and normalization methods without re-engineering the pipeline core. The observed efficacy of PCA-based intermediate integration suggests that latent-space representations are sufficient for preserving biological signal while mitigating the noise typically associated with high-throughput sequencing. Furthermore, the inclusion of an FHIR R4 export module addresses a critical gap in translational medicine, facilitating the direct injection of research findings into clinical decision support systems.

9.2 Limitations and Constraints

Despite its robust performance on TCGA-BRCA, MultiOmicsPipeline faces certain limitations. Memory consumption scales linearly with the number of features, which may necessitate a transition to distributed computing backends (e.g., Apache Spark) for datasets exceeding 10^5 samples. Additionally, while the PCA approach is computationally efficient, it can obscure the direct interpretability of individual gene contributions compared to more complex factor models like MOFA+ [2]. Future iterations will explore the integration of XAI (Explainable AI) modules to bridge this gap.

9.3 Development Roadmap

Our long-term roadmap focuses on enhancing the framework’s scalability and clinical utility. Table 9 details our planned milestones, including support for longitudinal clinical tracking and full OpenAPI-compliant REST services.

10 Conclusion

MultiOmicsPipeline represents a significant step towards the standardization of multi-modal data integration in precision oncology. By providing a structured, reproducible framework that bridges the gap between raw sequencing data and FHIR-compliant clinical exports, we empower researchers to focus on biological discovery rather than technical orchestration. Our validation on the TCGA-BRCA cohort demonstrates that the framework is not only computationally efficient but also analytical rigorous, achieving near-perfect data quality metrics. As we move towards a future of

Table 9: Projected development roadmap (2026–2027).

Milestone	Strategic Objective	Target KPI
Q2 2026	Methylation (850K) integration	>95% CpG coverage
Q3 2026	Dynamic D3.js Visualizer	Response <2 s
Q4 2026	OpenAPI 3.0 REST Services	99.9% Uptime
Q1 2027	Memory Optimization (Lazy Load)	−40% RAM footprint
Q3 2027	Apache Spark Distributed Engine	>100k samples
Q4 2027	FHIR R5 Semantic Alignment	Zero-error export

[23] Cancer Genome Atlas Network. Comprehensive molecular portraits of human breast tumours. *Nature*, 490(7418):61–70, 2012.

pan-cancer integration, this pipeline serves as a scalable foundation for the next generation of data-driven clinical diagnostics.

Acknowledgments

We acknowledge the TCGA Research Network and the NCI Genomic Data Commons for providing the public datasets used in this work. We also thank the Faculty of Sciences of Monastir for their institutional support. This project was conducted under the supervision of Prof. Mohsen Maraoui (Academic Year 2025–2026).

References

[1] K. Tomczak, P. Czerwińska, and M. Wiznerowicz. The Cancer Genome Atlas (TCGA): An immeasurable source of knowledge. *Contemporary Oncology*, 19(1A):A68–A77, 2015.

[2] R. Argelaguet et al. Multi-Omics Factor Analysis—a framework for unsupervised integration of multi-omics data sets. *Molecular Systems Biology*, 14(6):e8124, 2018.

[3] W.E. Johnson, C. Li, and A. Rabinovic. Adjusting batch effects in microarray expression data using empirical Bayes methods. *Biostatistics*, 8(1):118–127, 2007.

[4] M.E. Ritchie et al. limma powers differential expression analyses for RNA-sequencing and microarray studies. *Nucleic Acids Research*, 43(7):e47, 2015.

[5] I. Subramanian et al. Multi-omics data integration, interpretation, and its application. *Bioinformatics and Biology Insights*, 14:1–24, 2020.

[6] M. D. Wilkinson et al. The FAIR guiding principles for scientific data management and stewardship. *Scientific Data*, 3:160018, 2016.

[7] A. Colaprico et al. TCGAAbiolinks: an R/Bioconductor package for integrative analysis with GDC data. *Nucleic Acids Research*, 44(8):e71, 2016.

[8] HL7 International. HL7 FHIR Release 4 Specification, 2023. <https://hl7.org/fhir/R4/>.

[9] L. M. McShane et al. Criteria for the use of omics-based predictors in clinical trials. *BMC Medicine*, 11:220, 2013.

[10] D. B. Rubin. Inference and missing data. *Biometrika*, 63(3):581–592, 1976.

[11] O. Troyanskaya et al. Missing value estimation methods for DNA microarrays. *Bioinformatics*, 17(6):520–525, 2001.

[12] S. van Buuren and K. Groothuis-Oudshoorn. mice: Multivariate Imputation by Chained Equations in R. *Journal of Statistical Software*, 45(3):1–67, 2011.

[13] D.J. Stekhoven and P. Bühlmann. MissForest—non-parametric missing value imputation for mixed-type data. *Bioinformatics*, 28(1):112–118, 2012.

[14] M. I. Love, W. Huber, and S. Anders. Moderated estimation of fold change and dispersion for RNA-seq data with DESeq2. *Genome Biology*, 15(12):550, 2014.

[15] J. S. Parker et al. Supervised risk predictor of breast cancer based on intrinsic subtypes. *Journal of Clinical Oncology*, 27(8):1160–1167, 2009.

[16] F. Pedregosa et al. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.

[17] M. Rocklin. Dask: Parallel computation with blocked algorithms and task scheduling. In *Proceedings of the 14th Python in Science Conference*, pages 130–136, 2015.

[18] W. McKinney. Data structures for statistical computing in Python. In *Proceedings of the 9th Python in Science Conference*, pages 56–61, 2010.

[19] J.N. Weinstein et al. The Cancer Genome Atlas Pan-Cancer analysis project. *Nature Genetics*, 45(10):1113–1120, 2013.

[20] J. T. Leek et al. Tackling the widespread and critical impact of batch effects in high-throughput data. *Nature Reviews Genetics*, 11(10):733–739, 2010.

[21] J. C. Mandel et al. SMART on FHIR: a standards-based, interoperable apps platform for electronic health records. *Journal of the American Medical Informatics Association*, 23(5):899–908, 2016.

[22] I. T. Jolliffe. *Principal Component Analysis*. Springer, 2nd edition, 2002.